

CUSTOMER CHURN PREDICTION

Step 1: Import Libraries

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, confusion_matrix,
accuracy_score
```

Step 2: Load Dataset

```
df = pd.read_csv("WA_Fn-UseC_-Telco-Customer-Churn.csv") # Download
from Kaggle
print("Data loaded successfully!")
```

Data loaded successfully!

```
df.head(3)
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure
0	7590-VHVEG	Female	0	Yes	No	1
1	5575-GNVDE	Male	0	No	No	34
2	3668-QPYBK	Male	0	No	No	2

	MultipleLines	InternetService	OnlineSecurity	...
0	No phone service	DSL	No	...
1	No	DSL	Yes	...
2	No	DSL	Yes	...

	TechSupport	StreamingTV	StreamingMovies	Contract
0	PaperlessBilling			

0	No	No	No	Month-to-month
Yes				
1	No	No	No	One year
No				
2	No	No	No	Month-to-month
Yes				

	PaymentMethod	MonthlyCharges	TotalCharges	Churn
0	Electronic check	29.85	29.85	No
1	Mailed check	56.95	1889.5	No
2	Mailed check	53.85	108.15	Yes

[3 rows x 21 columns]

Step 3: Basic Exploration

```
print(df.shape)
print(df.columns)
print(df["Churn"].value_counts())

(7043, 21)
Index(['customerID', 'gender', 'SeniorCitizen', 'Partner',
      'Dependents',
      'tenure', 'PhoneService', 'MultipleLines', 'InternetService',
      'OnlineSecurity', 'OnlineBackup', 'DeviceProtection',
      'TechSupport',
      'StreamingTV', 'StreamingMovies', 'Contract',
      'PaperlessBilling',
      'PaymentMethod', 'MonthlyCharges', 'TotalCharges', 'Churn'],
      dtype='object')
Churn
No      5174
Yes     1869
Name: count, dtype: int64
```

Step 4: Handle Missing Values

```
print(df.isnull().sum())
df["TotalCharges"] = pd.to_numeric(df["TotalCharges"],
errors='coerce') # Convert to numeric
df["TotalCharges"].fillna(df["TotalCharges"].median(), inplace=True)

customerID      0
gender           0
SeniorCitizen   0
Partner         0
```

```
Dependents      0
tenure          0
PhoneService    0
MultipleLines   0
InternetService 0
OnlineSecurity  0
OnlineBackup    0
DeviceProtection 0
TechSupport     0
StreamingTV     0
StreamingMovies 0
Contract        0
PaperlessBilling 0
PaymentMethod   0
MonthlyCharges  0
TotalCharges    0
Churn           0
dtype: int64
```

C:\Users\Dell\AppData\Local\Temp\ipykernel_13684\4084226285.py:3:
FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df["TotalCharges"].fillna(df["TotalCharges"].median(), inplace=True)
```

Step 5: Drop Irrelevant Columns

```
df.drop(['customerID'], axis=1, inplace=True)
```

Step 6: Convert Categorical Variables to Numeric


```
categorical_cols = df.select_dtypes(include=['object']).columns
label_encoders = {}

for col in categorical_cols:
```

```
le = LabelEncoder()
df[col] = le.fit_transform(df[col])
label_encoders[col] = le
```

Step 7: Correlation Matrix

```
plt.figure(figsize=(12,8))
sns.heatmap(df.corr(), annot=True, cmap="coolwarm")
plt.title("Correlation Heatmap")
plt.show()
```



Step 8: Feature and Target Split

```
X = df.drop("Churn", axis=1)
y = df["Churn"]
```

Step 9: Train-Test Split

```
X_train, X_test, y_train, y_test = train_test_split(X, y,  
test_size=0.2, random_state=42)
```

Step 10: Train Model

```
model = RandomForestClassifier(n_estimators=100, random_state=42)  
model.fit(X_train, y_train)
```

```
RandomForestClassifier(random_state=42)
```

Step 11: Evaluate Model

```
y_pred = model.predict(X_test)  
print("Accuracy:", accuracy_score(y_test, y_pred))  
print("Classification Report:\n", classification_report(y_test,  
y_pred))  
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
```

Accuracy: 0.7955997161107168

Classification Report:

	precision	recall	f1-score	support
0	0.83	0.91	0.87	1036
1	0.66	0.47	0.55	373
accuracy			0.80	1409
macro avg	0.74	0.69	0.71	1409
weighted avg	0.78	0.80	0.78	1409

Confusion Matrix:

```
[[945  91]  
 [197 176]]
```

Step 12: Feature Importance

```
importances = pd.Series(model.feature_importances_, index=X.columns)  
importances.nlargest(10).plot(kind='barh', title="Top 10 Important  
Features")  
plt.xlabel("Feature Importance Score")  
plt.show()
```

Top 10 Important Features

